

Introduction to Deep Learning

Session 10: CNNs in Practice

May 2026

Magda Gregorová - magda.gregorova@thws.de



Licensed according to CC-BY-SA 4.0

Lecture Overview

Three Regimes

Regime	When	Data needed	Compute
Train from scratch	Novel architecture, very different domain	Large labelled dataset	Very high
Pretrain + fine-tune	Moderate dataset, research setting	Moderate	Medium
Download & adapt	Most practical tasks	Small-medium	Low

Modern workflow:

1. download pretrained from torchvision / timm / HuggingFace
2. adapt to task

Training from scratch rare - large organisations with compute and data budget

Why Pretrained Models Work

Key insight: CNN layers learn features of increasing specificity

Early layers - general:

- Edges, colours, orientations
- Useful across almost all visual tasks
- Transfer well to new domains

Later layers - specific:

- Object parts, textures, patterns
- Specific to source task
- Need adaptation for new task

Pretrained: ImageNet (1.2M images), LAION-5B (5.8B image-text pairs), proprietary

Works because: low-level visual features universal - edges exist in all images

Feature Extraction vs Fine-tuning

Feature extraction:

- Freeze all pretrained layers
 - Replace and train FC head only
- + Fast, few trainable parameters
+ Works with small datasets
+ No risk of catastrophic forgetting
- Limited adaptation to new domain

Fine-tuning:

- Unfreeze some or all layers
 - Train with small learning rate
- + Better performance
+ Adapts to new domain
- Needs more data
- Risk of catastrophic forgetting

Catastrophic forgetting: large learning rate overwrites pretrained weights - use $10\times$ – $100\times$ smaller learning rate than original training

When to Use Which

	Similar domain	Different domain
Small dataset	Feature extraction (freeze all, train head)	Feature extraction (try different layer depths)
Large dataset	Fine-tune all layers (small learning rate)	Fine-tune all layers (more epochs, careful tuning)

Rule of thumb:

- Less data → freeze more layers
- More similar to source domain → reuse more layers
- Always start with feature extraction, add fine-tuning if needed

Freezing strategy:

- Start fully frozen - train head only
- Unfreeze progressively from top (later layers first)
- Earlier layers = more general = freeze longer

Learning rates:

- New head: standard learning rate (e.g. $1e-3$)
- Fine-tuned layers: much smaller (e.g. $1e-5$)
- **Discriminative learning rates** - different rates per layer group
- Warmup - start small, increase gradually → new head stabilization

Other considerations:

- More aggressive data augmentation when dataset is small
- Early stopping essential - fine-tuning overfits quickly

Practical CNN Training

Two main sources of pretrained models:

torchvision.models:

- Official PyTorch model zoo
- ResNet, VGG, EfficientNet, ViT
- Pretrained on ImageNet
- Simple API

```
from torchvision.models import resnet50,  
    ResNet50_Weights  
model = resnet50(weights=ResNet50_Weights.DEFAULT)
```

timm (PyTorch Image Models):

- 700+ pretrained models
- State-of-the-art architectures
- Multiple pretraining datasets
- Community standard

```
import timm  
model = timm.create_model('resnet50',  
    pretrained=True)
```

HuggingFace: transformers library for vision-language models (CLIP, ViT, DINO)

Adapting Pretrained Model

Replace classifier head:

```
model = resnet50(weights=ResNet50_Weights.DEFAULT)

# freeze all layers
for param in model.parameters():
    param.requires_grad = False

# replace FC head
num_features = model.fc.in_features
model.fc = nn.Linear(num_features, num_classes)

# only head parameters are trainable
optimizer = torch.optim.Adam(model.fc.parameters(), lr=1e-3)
```

All basics hold:

```
model.train()
model.eval()
optimizer.zero_grad()
```

For fine-tuning: unfreeze layers selectively

```
# unfreeze last ResNet block
for param in model.layer4.parameters():
    param.requires_grad = True

# fine-tune with discriminative learning rate
optimizer = torch.optim.Adam([
    {'params': model.layer4.parameters(), 'lr': 1e-5},
    {'params': model.fc.parameters(), 'lr': 1e-3},
])
```

Standard training/ validation / test pipeline:

```
from torchvision import transforms

train_transform = transforms.Compose([
    transforms.RandomResizedCrop(224),
    transforms.RandomHorizontalFlip(),
    transforms.ColorJitter(brightness=0.2, contrast=0.2),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])

val_transform = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(...)
])
```


- No augmentation at validation/ test - deterministic transforms only
- Normalisation statistics must match pretraining - check torchvision / timm

Beyond Classification


Vision Tasks




Classification



Object detection




Semantic segmentation



Instance segmentation

All built on CNN backbones - pretrained feature extractors + task-specific heads

Object Detection




Output: bounding boxes + class labels + confidence scores per object instance

Annotation: bounding box around every object instance

Two main strategies:

- **Two-stage** (**Faster R-CNN**): region proposal network → classify and refine - higher accuracy, slower
- **One-stage** (**YOLO**, **SSD**): directly predict boxes and classes in single pass - faster, real-time capable

Evaluation: IoU and mAP - see next slide




IoU (Intersection over Union):

- Overlap between predicted and ground truth box
- Prediction correct if $\text{IoU} > \text{threshold}$ (e.g. 0.5)
- $\text{IoU} = \frac{|\text{intersection}|}{|\text{union}|} \in [0, 1]$

Average Precision (AP):

- Vary confidence threshold - plot precision-recall curve
- AP = area under that curve per class
- mAP = mean AP across all classes

Semantic Segmentation



Output: pixel-wise class label map - same spatial size as input

Annotation: careful pixel-level boundary tracing


Main strategy: encoder-decoder architecture

- **Encoder:** CNN backbone reduces resolution, extracts features
- **Decoder:** upsamples back to input resolution
- **Skip connections:** recover spatial detail lost in encoding

Key methods: U-Net, DeepLab v3, SegFormer

Evaluation: mean Intersection over Union (mIoU) - mean IoU across all classes

Instance Segmentation



Output: per-instance binary mask + class label - distinguishes individual objects of same class

Annotation: per-instance pixel masks

Main strategy: detect then segment

- **Mask R-CNN:** extends Faster R-CNN with mask prediction head - binary mask per detected box
- **SAM** (Segment Anything): foundation model, prompt-based - any object, zero-shot

Panoptic segmentation: unified - semantic for background, instance for objects - modern standard

Evaluation: mAP on masks rather than boxes

Task	Annotation type	Time per image	Example dataset
Classification	Image-level label	Seconds	ImageNet
Detection	Bounding boxes	Minutes	COCO
Semantic segmentation	Pixel-wise labels	30–90 min	Cityscapes
Instance segmentation	Per-instance masks	Hours	COCO

Annotation cost = the real bottleneck - not model complexity

Modern tools to reduce cost:

- **SAM** (Segment Anything Model) - interactive, prompt-based annotation assistance
- Semi-supervised and weakly supervised methods - learn from cheap labels
- Synthetic data - generate annotated data automatically

Transfer learning:

- Download pretrained model
- Feature extraction vs fine-tuning
- Freeze progressively

Practical tools:

- Pretrained model libraries
- Image transformations

Beyond classification:

- Object detection
- Semantic segmentation
- Instance segmentation
- All use CNN backbones