

Outline

1. What is Deep Learning?

2. From Linear Models to Neural Networks

3. Neural Network Architecture

4. Loss Functions and Learning

5. Course Overview

Outline

1. What is Deep Learning?

2. From Linear Models to Neural Networks

3. Neural Network Architecture

4. Loss Functions and Learning

5. Course Overview

Deep Learning: A Brief History

The Journey from Perceptrons to Modern AI

- ▶ **1958:** Rosenblatt's Perceptron - single layer, limited capacity
- ▶ **1969:** Minsky & Papert show perceptron limitations (e.g., XOR problem)
- ▶ **1986:** Backpropagation algorithm popularized (Rumelhart et al.)
- ▶ **1990s-2000s:** "AI Winter" - neural networks fall out of favor
- ▶ **2012:** AlexNet wins ImageNet - deep learning renaissance begins
- ▶ **2017-present:** Transformers revolutionize NLP and beyond

Why now?

- ▶ Massive datasets (ImageNet, Common Crawl, etc.)
- ▶ Computational power (GPUs, TPUs)
- ▶ Algorithmic innovations (ReLU, batch normalization, residual connections)

What Can Deep Learning Do?

Computer Vision

- ▶ Image classification, object detection, semantic segmentation
- ▶ Face recognition, medical image analysis

Natural Language Processing

- ▶ Machine translation, text generation, question answering
- ▶ Large language models (GPT, Claude, Gemini)

Other Domains

- ▶ Speech recognition and synthesis
- ▶ Game playing (AlphaGo, AlphaZero)
- ▶ Protein structure prediction (AlphaFold)
- ▶ Autonomous driving, robotics

Common thread: Learning hierarchical representations from data

Outline

1. What is Deep Learning?

2. From Linear Models to Neural Networks

3. Neural Network Architecture

4. Loss Functions and Learning

5. Course Overview

Linear Models: A Quick Recap

Linear Regression

- ▶ Model: $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$
- ▶ Input: $\mathbf{x} \in \mathbb{R}^d$, weights: $\mathbf{w} \in \mathbb{R}^d$, bias: $b \in \mathbb{R}$
- ▶ Learns a hyperplane in input space

Linear Classification (e.g., Logistic Regression)

- ▶ Model: $f(\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + b)$ where $\sigma(z) = \frac{1}{1+e^{-z}}$
- ▶ Decision boundary is still linear

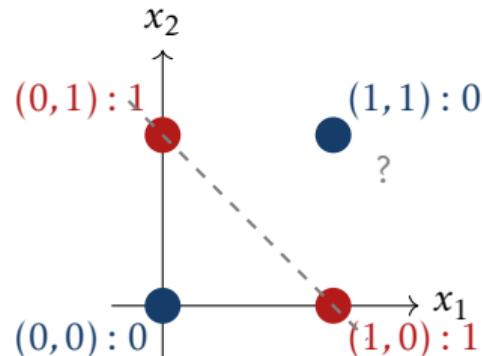
Fundamental Limitation

- ▶ Can only learn linear decision boundaries
- ▶ Many real-world problems are inherently non-linear
- ▶ Example: XOR problem cannot be solved by linear classifier

The XOR Problem

Problem Definition

- XOR (exclusive OR): output is 1 if inputs differ, 0 otherwise



Why Linear Models Fail

- No single line can separate blue points from red points
- We need non-linear decision boundaries
- Solution: multiple layers with non-linear activations

Outline

1. What is Deep Learning?

2. From Linear Models to Neural Networks

3. Neural Network Architecture

4. Loss Functions and Learning

5. Course Overview

Biological Inspiration

- ▶ Loosely inspired by biological neurons
- ▶ Multiple inputs, single output
- ▶ Non-linear activation function

Mathematical Model

- ▶ Input: $\mathbf{x} = (x_1, x_2, \dots, x_d)$
- ▶ Weights: $\mathbf{w} = (w_1, w_2, \dots, w_d)$, bias: b
- ▶ Pre-activation: $z = \sum_{i=1}^d w_i x_i + b = \mathbf{w}^T \mathbf{x} + b$
- ▶ Output: $a = \phi(z)$ where ϕ is the activation function

Key insight: Without ϕ , composition of neurons is still just linear!

Activation Functions

Common Choices

Sigmoid: $\sigma(z) = \frac{1}{1+e^{-z}}$

- ▶ Output range: $(0, 1)$
- ▶ Historically popular, now less common in hidden layers
- ▶ Problem: vanishing gradients for $|z|$ large

Tanh: $\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$

- ▶ Output range: $(-1, 1)$
- ▶ Zero-centered (better than sigmoid)
- ▶ Still suffers from vanishing gradients

ReLU (Rectified Linear Unit): $\text{ReLU}(z) = \max(0, z)$

- ▶ Most popular for hidden layers
- ▶ Advantages: fast to compute, no vanishing gradient for $z > 0$
- ▶ Disadvantage: “dead neurons” when $z \leq 0$ always

Multi-Layer Perceptron (MLP)

Layer-wise Architecture

- ▶ **Input layer:** raw features $\mathbf{x} \in \mathbb{R}^{d_0}$
- ▶ **Hidden layers:** L layers with dimensions d_1, d_2, \dots, d_L
- ▶ **Output layer:** predictions $\hat{\mathbf{y}} \in \mathbb{R}^{d_{L+1}}$

Forward Pass for Layer ℓ

$$\mathbf{z}^{(\ell)} = \mathbf{W}^{(\ell)} \mathbf{a}^{(\ell-1)} + \mathbf{b}^{(\ell)}$$

$$\mathbf{a}^{(\ell)} = \phi(\mathbf{z}^{(\ell)})$$

where:

- ▶ $\mathbf{W}^{(\ell)} \in \mathbb{R}^{d_\ell \times d_{\ell-1}}$ is the weight matrix
- ▶ $\mathbf{b}^{(\ell)} \in \mathbb{R}^{d_\ell}$ is the bias vector
- ▶ ϕ is applied element-wise
- ▶ Convention: $\mathbf{a}^{(0)} = \mathbf{x}$ (input)

Network Depth and Width

Network Capacity

- ▶ **Width:** number of neurons per layer
- ▶ **Depth:** number of hidden layers
- ▶ Both affect the model's representational power

Universal Approximation Theorem (informal)

- ▶ A neural network with **one hidden layer** of sufficient width can approximate any continuous function on a compact domain
- ▶ **But:** may require exponentially many neurons!
- ▶ **Depth is more parameter-efficient than width**

Why Deep Networks?

- ▶ Learn hierarchical representations (edges → shapes → objects)
- ▶ More parameter-efficient for complex functions
- ▶ Empirically perform better on real tasks

Outline

1. What is Deep Learning?

2. From Linear Models to Neural Networks

3. Neural Network Architecture

4. Loss Functions and Learning

5. Course Overview

The Learning Problem

Supervised Learning Setup

- ▶ Training data: $\mathcal{D} = \{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^N$
- ▶ Goal: find parameters $\boldsymbol{\theta} = \{\mathbf{W}^{(1)}, \mathbf{b}^{(1)}, \dots, \mathbf{W}^{(L+1)}, \mathbf{b}^{(L+1)}\}$ that minimize prediction error

Loss Function

- ▶ Measures how well the network fits the data
- ▶ For a single example: $\mathcal{L}(\hat{\mathbf{y}}, \mathbf{y})$
- ▶ Total loss (empirical risk):

$$\mathcal{L}(\boldsymbol{\theta}) = \frac{1}{N} \sum_{i=1}^N \mathcal{L}(f(\mathbf{x}_i; \boldsymbol{\theta}), \mathbf{y}_i)$$

Optimization Goal

$$\boldsymbol{\theta}^* = \arg \min_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta})$$

Common Loss Functions

Regression: Mean Squared Error (MSE)

$$\mathcal{L}_{\text{MSE}}(\hat{\mathbf{y}}, \mathbf{y}) = \frac{1}{2} \|\hat{\mathbf{y}} - \mathbf{y}\|^2 = \frac{1}{2} \sum_{j=1}^d (\hat{y}_j - y_j)^2$$

Probabilistic interpretation: assumes $y \sim \mathcal{N}(\hat{y}, \sigma^2)$, equivalent to maximum likelihood

Binary Classification: Binary Cross-Entropy

$$\mathcal{L}_{\text{BCE}}(\hat{y}, y) = -y \log(\hat{y}) - (1 - y) \log(1 - \hat{y})$$

where $\hat{y} = \sigma(\mathbf{w}^T \mathbf{x} + b) \in (0, 1)$ is predicted probability

Probabilistic interpretation: negative log-likelihood for Bernoulli distribution

Multi-Class Classification: Cross-Entropy

Softmax Activation

$$\hat{y}_j = \frac{e^{z_j}}{\sum_{k=1}^C e^{z_k}} \quad \text{for } j = 1, \dots, C$$

- ▶ Converts logits \mathbf{z} to probability distribution
- ▶ $\sum_{j=1}^C \hat{y}_j = 1$ and $\hat{y}_j \in (0, 1)$

Cross-Entropy Loss

$$\mathcal{L}_{\text{CE}}(\hat{\mathbf{y}}, \mathbf{y}) = - \sum_{j=1}^C y_j \log(\hat{y}_j)$$

For one-hot encoded labels (e.g., $\mathbf{y} = [0, 1, 0, 0]$ for class 2):

$$\mathcal{L}_{\text{CE}} = -\log(\hat{y}_c)$$

where c is the true class

Probabilistic interpretation: negative log-likelihood for categorical distribution

Gradient Descent Preview

The Optimization Strategy

- We need to minimize $\mathcal{L}(\theta)$ over all parameters
- Problem: high-dimensional, non-convex, no closed-form solution
- Solution: **iterative gradient-based optimization**

Gradient Descent (simplified)

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} \mathcal{L}(\theta_t)$$

- $\eta > 0$ is the learning rate (step size)
- Move in direction opposite to gradient (steepest descent)
- Repeat until convergence

Key Challenge How do we compute $\nabla_{\theta} \mathcal{L}$ efficiently for deep networks?
Answer: Backpropagation (next lecture!)

Outline

1. What is Deep Learning?

2. From Linear Models to Neural Networks

3. Neural Network Architecture

4. Loss Functions and Learning

5. Course Overview

Course Structure

Weeks 1-3: Foundations

- ▶ Neural network basics, backpropagation, computational graphs
- ▶ Optimization algorithms (SGD, momentum, Adam)
- ▶ Regularization techniques (dropout, batch normalization, weight decay)

Weeks 4-6: Convolutional Networks

- ▶ CNN architectures and design principles
- ▶ Transfer learning and data augmentation

Weeks 7-9: Sequential Models

- ▶ RNNs, LSTMs, GRUs
- ▶ Attention mechanisms and Transformers

Weeks 10-14: Advanced Topics

- ▶ Generative models, self-supervised learning
- ▶ Practical considerations: training at scale, deployment
- ▶ Ethics and fairness

Course Components

- ▶ **Lectures:** 1.5 hours per week (theory and concepts)
- ▶ **Exercises:** 1.5 hours per week (implementation and problem-solving)
- ▶ **Semester Project:** Train models on real datasets
- ▶ **Parallel course:** Basic ML (covers general ML foundations)

Prerequisites

- ▶ Programming: Python, NumPy (we'll use PyTorch or TensorFlow)
- ▶ Math: Linear algebra, calculus, basic probability
- ▶ We'll provide probability refreshers just-in-time as needed

Recommended Reading

- ▶ Simon J.D. Prince: *Understanding Deep Learning* (2023)
- ▶ Free online: <https://udlbook.github.io/udlbook/>

Next Week: Backpropagation and Optimization

What We'll Cover

- ▶ Computational graphs and automatic differentiation
- ▶ Backpropagation algorithm in detail
- ▶ Stochastic gradient descent and mini-batches
- ▶ Common optimization algorithms (momentum, Adam)

Preparation

- ▶ Review: chain rule from calculus
- ▶ Review: matrix calculus (gradients, Jacobians)
- ▶ Reading: Prince Chapter 6 (Fitting Models), Chapter 7 (Gradients and Initialization)

Exercise Session This Week

- ▶ Python/NumPy refresher
- ▶ Implementing activation functions
- ▶ Forward propagation from scratch